MapReduce Tutorial

I. Introduction.

Why MapReduce?

Traditional Enterprise System

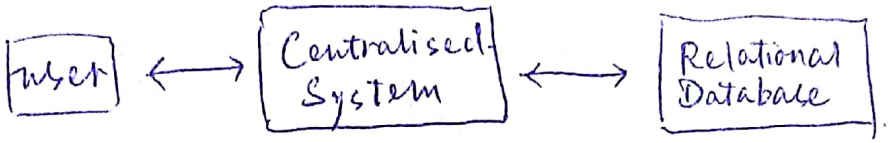

Limitation: Centralised system creates too much of a bottleneck while processing multiple files simultaneously.

Solution: MapReduce Algorithm.

MR. divides a task into small parts and assign them to many computers. Results are collected at one place and integrated to form the result dataset.

How MapReduce works?

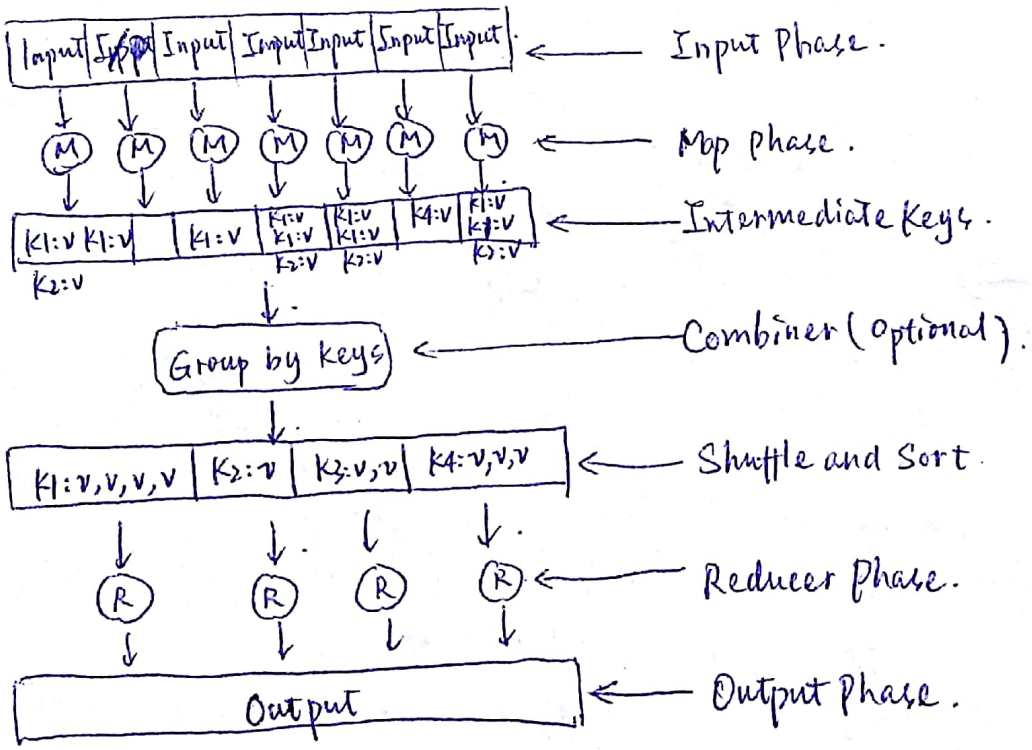

Input Phase: Record Reader. parses data to the mapper in <Key, value>

Map: a user-defined function, <key, value> → new <key, value>

Combiner: In a mapper, combiner is a local reducer.

<intermediate key, value> → combiner does aggregation.

Shuffle and Sort: Reducer task starts here.

Grouped <k, v> pairs → local machine (Reducer here).

<k, v> sorted by key → data list.

Reducer : grouped key-value paired data ~~as input~~.
↓
Reducer function on each of them.
aggregated, filtered, combined.
↓
new <k, v> pairs.

Output phase: write <k, v> pairs onto a file.

Example:

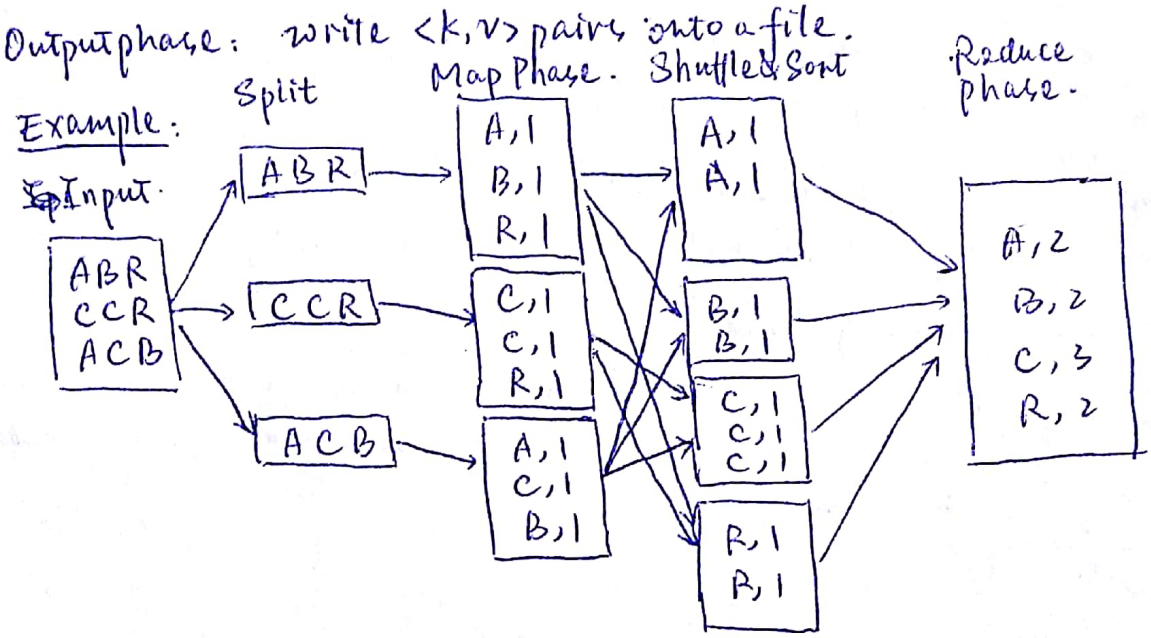

- MapReduce Example (twitter).

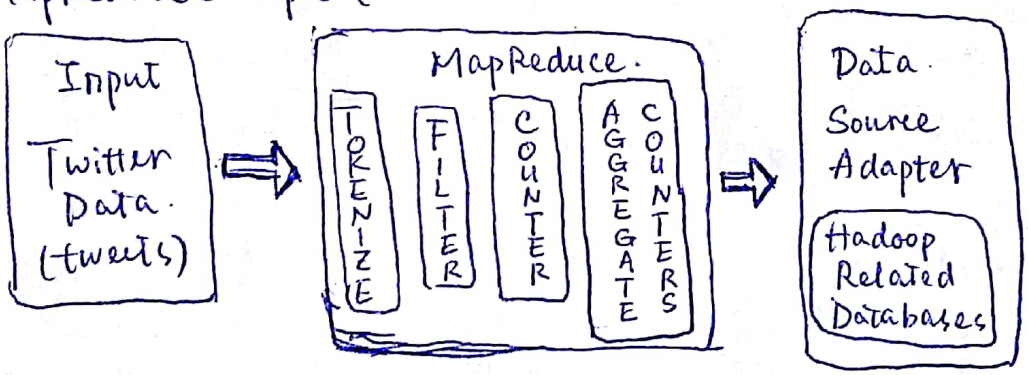

Tokenize : tokenizes tweets into maps of tokens. <k, v> pairs
Filter : filters unwanted words, writes filtered words → <k, v>.
Count : generates a token counter per word.
Aggregate counters : prepare an aggregate of similar counter values.

MapReduce Tutorial

II. Algorithm

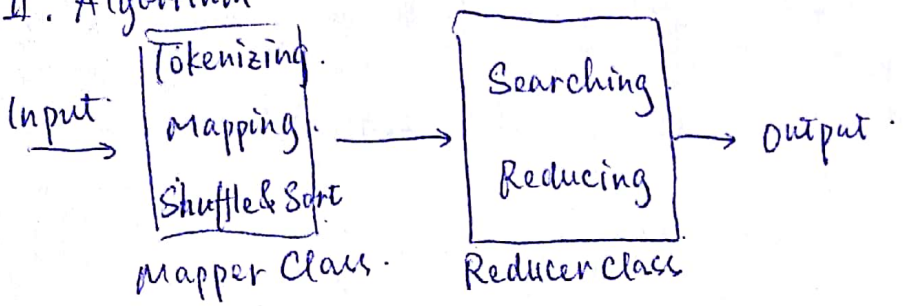

1. Sorting: output of mapper class by their keys.
   <K,V>
2. Searching: combiner phase & reducer phase.
3. Indexing: inverted index, TF-IDF. (Term freq. – inverted doc. freq.).

$$TF_{(the)} = \frac{\#\text{ of 'the' terms in the doc.}}{\#\text{ of terms in the doc.}}$$

$$IDF(the) = \log \frac{\text{total \# of documents}}{\text{e. \# of documents with term 'the'}}$$

III Installation (skip).

IV. API.

1. JobContext Interface:
   - super interface for all classes.
   - defines different jobs in mapreduce.

   Sub interfaces:
   - MapContext < KEYIN, VALUEIN, KEYOUT, VALUEOUT >.
   - ReduceContext < . . . >.

2. Job Class:
   - most important API in MapReduce.
   - configure the job, submit it, control its execution, query the state.
3. ~~Constructors of Job Class~~.
3. Mapper Class:
   - defines map job. input <key, value> ⟶ intermediate <key, value>
4. Reducer Class:
   - defines reduce job.
   - 3 primary phases
     (I) Shuffle: copies the sorted output from each mapper using HTTP across the network.
     (II) Sort: merge sorts reducer inputs by keys.
     Shuffle and sort phases occur simultaneously.

(3). Reduce: In this phase, the reduce(object, iterable, context) method is called for each <key, (collection of values)> in the sorted inputs.
- reduce method is called once for each key value.

V. Hadoop Implementation.

1. MapReduce Algorithm.
- Hadoop sends Map and Reduce tasks to appropriate servers in the cluster
- Most computing take place on nodes with data locally.
- after computation completes, collect to form appropriate result and sends it back to the Hadoop server.

2. Input and Outputs.
- <K, V> pairs as input and output.
- Serializable → Writable interface.
- key classes implement the WritableComparable interface to facilitate sorting.

Input → map → Intermediate <K2, V2> → reduce → output: <K3, V3>
<K1, V1>

3. MapReduce Implementation. (Skip).

VI. Partitioner.

Properties:
- takes place after Map phase and before Reduce phase.
- # partitioner = # reducer, determine # partitioner according to # reducer.
- partitions intermediate <K, V> from mapper.
- works like a hash fn.

VII. Combiners.

Properties:
(1) Semi-reducer.
(2) Summarize map output records with the same key.
(3)* reduce the volume of data transfer between map and reduce.
Note: usually map output volume is high.

How combiner works?
(1) no predefined interface. implement the reducer interface's reduce() method.
(2) produce summary info. from a large dataset because it replaces the original Map out.

MapReduce Tutorial

VII. Combiner (con't).

How combiner works?

(3) usually the code and operation for a combiner is similar to that of a reducer.

*(4) The combiner phase takes each key-value pair from the Map phase, processes it, and produces the output as key-value collection pairs.

e.g. wordcount.java.